

# Data Taming R reminder sheet

Bill S. Preston Esq.

2024-08-01

## 1 Setup

### 1.1 Initialising knitr

- This first code chunk will execute, but not display in the output because of the option `include=FALSE` in the definition of the code chunk.

knitr general options:

- `message=FALSE` when setting the `knitr` general options to suppress messages when loading the packages later in the file.
- `results=FALSE` to suppress output. (Change this to `results=TRUE` to see the output.)

### 1.2 Loading packages

```
library(tidyverse)
library(inspectdf)
library(caret)
library(moments)
library(tidymodels)
library(modelr)
library(ISLR)
library(car)
```

## 2 Loading data

- Use `<-` to assign new datasets and variables

### 2.1 Data already in the library

```
data("mpg")
mpg1<-mpg
```

### 2.2 Read in a csv file

```
pop1<- read_csv("population.csv")
```

## 3 Displaying data

### 3.1 Displaying data frames and tibbles

```
mpg1 # Prints the first 10 lines
print(n=10, mpg) # Also prints the first 10 lines
head(mpg1, 9) # Also prints the first 9 lines
tail(mpg1, 8) # Prints the final 8 lines
View(mpg1) # This command will display data in a new window
```

### 3.2 To show the number of rows and columns of the data

- `dim()` gives a list of the form `[\# rows, \# columns]`

```
dim(mpg)
```

- To just find the number of rows use `nrow()`

```
nrow(mpg)
```

- To just find the number of columns use `ncol()`

```
ncol(mpg)
```

### 3.3 To access a particular column

```
mpg1$cyl
```

## 4 Extracting data

### 4.1 Extracting columns

```
select(mpg, cyl:fl)
```

### 4.2 Extracting specific row numbers

```
mpg[5,]
```

### 4.3 Extracting specific values

#### 4.3.1 Just keeping the value

```
mpg[5,]$trans
class(mpg[5,]$trans)
```

#### 4.3.2 Putting the value into its own new tibble

```
mpg[5, "trans"]
class(mpg[5, "trans"])
```

#### 4.4 Extracting rows that match a TRUE/FALSE condition

```
filter(mpg, displ==3.1)
filter(mpg, between(mpg$displ,2.8,3.1))
```

## 5 Missing data

### 5.1 To check if there are any missing values in any columns

```
inspect_na(starwars)
```

### 5.2 Finding missing values in a specific column

- This returns a boolean list of TRUE/FALSE indicating the rows with missing data.
- This can be combined with the filter() command

```
is.na(starwars$species)
```

## 6 Manipulating data

### 6.1 Sorting a column

```
arrange(mpg, displ)
```

### 6.2 mutate() To add, change or remove columns

- Add column to right of dataset
  - mutate(dataset, new\_column\_name = value)

```
mutate(mpg1, IDnum=c(1:234))
mutate(mpg1, cty_hwy_avg=(cty -hwy)/2)
```

- Delete a column
  - mutate(dataset, existing\_column\_name = NULL)

```
mutate(mpg1, model=NULL)
```

- Change a column
  - mutate(dataset, existing\_column\_name = value)

```
mutate(mpg, displ=displ*10)
```

### 6.3 rename() to rename a column

- Use the syntax rename(dataset, new\_column\_name=old\_column\_name)

```
rename(mpg, displacement=displ)
```

### 6.4 relocate() to move a column

- Move a column “before” (to the left) of another column

```
relocate(mpg, "cyl", .before = model)
```

- Move a column “after” (to the right) of another column

```
relocate(mpg, "cyl", .after = cty)
```

## 6.5 To concatenate tibbles (glue them together)

- Use `bind_cols()`.

```
bind_cols(  
  mpg["trans"],  
  mpg["cty"]  
)
```

- Note that everything you are binding together should be a tibble, with unique variable names, otherwise the variable names in the resulting tibble will be nonsense.

```
bind_cols(  
  1:50,  
  1:50  
)
```

## 6.6 Grouping rows

- Group a set of rows together based on the values in one of the columns. Eg. this will group all the cars together by their number of cylinders.
- This can be used with the `summarise()` command to compute statistics for each group.

```
group_by(mpg, cyl)
```

## 6.7 Change data types

- Convert to factor (nominal categorical variable)

```
mpg1$cyl<-as.factor(mpg1$cyl)
```

- Convert to ordered factor (ordinal categorical variable)

```
mpg1$cyl<-as.ordered(mpg1$cyl)  
#This next bit of code will change the order of the levels  
mpg1$cyl<-factor(mpg1$cyl, levels=c("5", "8", "4", "6"), ordered=TRUE)
```

- Convert to numerical variable

```
mpg1$cyl<-as.numeric(mpg1$cyl)
```

- Convert to integer variable

```
mpg1$cyl<-as.integer(mpg1$cyl)
```

- Convert to character string

```
mpg1$cyl<-as.character(mpg1$cyl)
```

- Convert to logical/Boolean variable

– First need a column of TRUE/FALSE or 1/0

```
mpg2<-mutate(mpg1, tf="TRUE")  
mpg2$tf<-as.logical(mpg2$tf)
```

- Convert to date object

- Using lubridate package commands ymd or dmy

```
mpg2<-mutate(mpg1,date="2025-10-05")
mpg2$date<-ymd(mpg2$date)
mpg2<-mutate(mpg1,date="05-10-2025")
mpg2$date<-dmy(mpg2$date)
```

## 6.8 Rename entries in column

- Using fct\_recode(mpg\$drv, "new1"="old1", "new2"="old2", "new3"="old3")
- Only works when column is a factor or character string
- Also converts column to factor type

```
fct_recode(mpg$drv, "front"="f", "4x4"="4", "rear"="r")
```

# 7 Control structures

## 7.1 Decisions

The ifelse command has syntax: ifelse(condition, return if true, return if false)

```
ifelse(10==0,1,0)
ifelse(mpg$cyl==4,mpg$cyl,-99)
```

This can be combined with mutate() to selectively modify tibbles:

```
mutate(mpg, take4=ifelse(mpg$cyl==4,mpg$cyl,-99))
```

## 7.2 Loops

- R is a vectorised programming language, and so it is not optimised for loops. Therefore we do not use loops in this course, and you must find another way to achieve your goal.

# 8 A sequence of numbers

```
1:50
```

## 8.1 A sequence of numbers with step size

To define a sequence from  $a$  to  $b$  with steps of size  $s$ , use: seq(a,b,s)

```
seq(5,8,0.05)
```

# 9 Random sampling

## 9.1 Random sampling from a tibble

- Use sample\_n()
- Uniformly at random choose 20 rows with replacement

```
sample_n(mpg, 20, replace=TRUE)
```

- Uniformly at random choose 20 rows without replacement

```
sample_n(mpg, 20, replace=FALSE)
```

## 9.2 Generating a list of random numbers

- Use `sample()`
- Generate 7 random integers from 1 to 50 with replacement

```
sample(1:50, 7, replace = TRUE)
```

- Generate 7 random integers from 1 to 50 without replacement

```
sample(1:50, 7, replace = FALSE)
```

## 9.3 Setting the seed for a random number generation

```
set.seed(1234)
```

# 10 Character string manipulation

## 10.1 Special characters

When using regular expressions you need the following commands for special characters:

- `$`: use `\\$`
- `(`: use `\\(`

## 10.2 Joining (concatenating) strings

- These commands `paste0` and `str_c` seem to do the same thing

```
middle<-"middle bit;"
paste0("first bit;",middle, " last bit")
str_c("first bit;", middle, " last bit")
```

## 10.3 Extracting numbers from strings

- To extract numbers from strings use regular expressions. Eg. `(\\d+)`

```
df <- tibble(
  treatment = c("A", "B", "C"),
  response = c(12, 11, 10),
  some_text1 = c("abc 7", "abc 2", "abc 5"),
  some_text2 = c("abc 7 xyz 9", "abc 1 xyz 21", "abc 0 xyz 2"),
  some_text3 = c("abc 7 xyz 9", "abc -2 xyz 21", "abc 0.5 xyz 2")
)
str_match(df$some_text1, "abc (\\d+)")
```

- Also works for extracting multiple values from a string

```
str_match(df$some_text2, "abc (\\d+) xyz (\\d+)")
```

- Only works for positive integers

```
str_match(df$some_text3, "abc (\\d+) xyz (\\d+)")
```

## 10.4 Extracting alphabetic characters from strings

- Works for both upper and lower case letters

```
str_match(df$some_text2, "x([:alpha:]+)")
```

## 10.5 Replacing parts of strings

```
mpg[6,]$trans  
str_replace(mpg[6,]$trans, "a", "X")  
str_replace(mpg[6,]$trans, "\\(", "X")
```

### 10.5.1 To replace all parts of the string matching the pattern

```
mpg[6,]$trans  
str_replace_all(mpg[6,]$trans, "a", "X")
```

## 11 Precision

- Rounding off to  $n$  decimal places
  - Note that  $n$  can be zero or negative. (Experiment with it to see what it does.)

```
round(15.32257, 3)  
round(15.32257, -1)
```

- Rounding off to  $n$  significant figures

```
signif(15.32257, 3)
```

## 12 Statistics

- Count the number of rows that match each value of one of the columns

```
count(mpg, displ)  
count(mpg, drv)
```

- Maximum and minimum values of a numeric list

```
max(mpg$displ)  
min(mpg$displ)  
  
#Be careful: any NA values will mean NA is returned  
max(starwars$height)  
#To ignore the NA values use na.rm=TRUE  
max(starwars$height, na.rm=TRUE)
```

- Mean

```
mean(mpg$hwy)
```

- Sample standard deviation. (Note that the sample deviation uses  $N - 1$  in the denominator of the calculation.)

```
sd(mpg$hwy)
```

- Skewness

Use the `moments` package command, as some algorithms produce different results.

```
moments::skewness(mpg$hwy)
```

- The inter-quartile range

```
IQR(mpg$hwy)
```

## 12.1 Statistics on grouped data

- We can use the statistics commands via the `summarise()` command, which is especially useful for working with grouped data

```
summarise(mpg, mean_hwy = mean(hwy))  
summarise(group_by(mpg, cyl), mean_hwy = mean(hwy))
```

## 12.2 Summary statistics

- Calculate summary statistics for all numerical variables we can use `inspect_num()`

```
inspect_num(mpg)
```

## 12.3 Building formulae

- To write a formula we put the response variable on the left of `~` and the predictors on the right. Eg. `y ~ x + z`.
- To include all variables as predictors (except the response `y`) use a full-stop.
  - For example, if we have a set of predictors `x1`, `x2`, `x3` then `y ~ .` is equivalent to `y ~ x1 + x2 + x3`
- We can include interactions between our predictors with the colon `:`. Eg. `y ~ x + z + x:z`
- The easy way to write a formula with all individual terms and second-order interactions as predictors is `y ~ .^2`.
  - For example, if we have a set of predictors `x1`, `x2`, `x3` then `y ~ .^2` is equivalent to `y ~ x1 + x2 + x3 + x1:x2 + x1:x3 + x2:x3`

## 12.4 Linear models

```
df<- tibble(  
  x=c(1:20),  
  y=x+rnorm(20,0,2)  
)  
df_lm<-lm(y~x,df)  
summary(df_lm)
```


- Extracting the model coefficients

```
as.numeric(df_lm$coefficients)
```


### 12.4.1 Graphs for checking assumptions of linear models

```
plot(df_lm, which=1)
```






```
plot(df_lm, which=3)
```



```
plot(df_lm, which=2)
```



#### 12.4.2 Predicting with linear models

(Note that the explanatory values used for the prediction must be stored in a tibble/dataframe.)

- Prediction with prediction interval of level 85%

```
pred_values<-tibble(
  x=c(2.5, 7.2)
)
predict(df_lm, pred_values, interval="prediction", level = 0.85)
```

- Prediction with confidence interval of level 99%

```
predict(df_lm, pred_values, interval="confidence", level = 0.99)
```

### 12.5 Logistic models

- We first make a binary categorical variable in a data set

```
car_seats <- as_tibble(Carseats)
car_seats
car_seats <- car_seats %>%
  mutate("sales_high"=ifelse(Sales>8,"high","low"), .after = Sales)
car_seats$sales_high <- factor(car_seats$sales_high)
car_seats_1 <- car_seats %>%
  mutate(Sales=NULL)
```

### 12.5.1 Building the model

```
classification_lr <- logistic_reg() %>%  
  set_engine("glm")  
lrfit <- classification_lr %>%  
  fit(sales_high ~ Price, data = car_seats_1)
```

### 12.5.2 Predicting with the logistic model

```
predict(lrfit, new_data=car_seats_1)  
predict(lrfit, new_data=car_seats_1, type="prob")
```

## 12.6 Extracting the model data for general linear models

- Summary of the model

```
summary(lrfit$fit)
```

- Just the coefficients

```
as.numeric(lrfit$fit$coefficients)
```

## 12.7 Analysis of variance

```
Anova(lrfit$fit)
```

## 12.8 Building models with categorical predictors

- To see what new variables are introduced we can use `model_matrix()`.
- R will introduce a new binary variable for each level of the categorical predictor, except the reference level. The name of the new variable will be the concatenation of “variable name” and “level name”.

```
model_matrix(mpg, ~drv)
```

- Note that the data type must not be quantitative (`<int>` or `<dbl>`), otherwise the levels won’t be assigned to new variables.

```
model_matrix(mpg, ~cyl)
```

## 13 Testing and training sets

### 13.1 Splitting the data

- You can set the proportion of rows in the training set with the option `prop = ...`

```
mpg_split<-initial_split(mpg, prop=0.7)
```

### 13.2 Making training and testing sets

```
mpg_train <- training(mpg_split)  
mpg_test  <- testing(mpg_split)
```

## 14 Evaluating models

### 14.1 Regression models

- First we build a model on the training set

```
mpg_lm1 <- lm(hwy ~ cty, data = mpg_train)
```

- Then we predict on the testing set and put the predictions together with the true values

```
predict(mpg_lm1, mpg_test)
rpreds <- tibble(
  reg_truth = mpg_test$hwy,
  reg_preds = predict(mpg_lm1, mpg_test)
)
```

- To evaluate the model we use the `metrics()` command

```
metrics(rpreds, reg_preds, truth = reg_truth)
```

### 14.2 Classification models

- We train the model on the training set

```
car_split <- initial_split(car_seats_1)
car_train <- training(car_split)
car_test <- testing(car_split)

lrfit1 <- classification_lr %>%
  fit(sales_high ~ Price, data = car_train)
```

- Then we predict on the testing set, and collect the predictions together

```
logpreds <- bind_cols(
  car_test["sales_high"],
  predict(lrfit1, new_data = car_test, type = "class"),
  predict(lrfit1, new_data = car_test, type = "prob")
)
```

#### 14.2.1 Confusion matrix

```
cm <- logpreds %>%
  conf_mat(
    .pred_class,
    truth = sales_high
  )
```

#### 14.2.2 Accuracy

```
logpreds %>% accuracy(
  .pred_class,
  truth = sales_high
)
```

### 14.2.3 Sensitivity

- Note that you may have to change `event_level` between “first” and “second” depending on what counts as “success” in your model.

```
logpreds %>% sens(  
  .pred_class,  
  truth = sales_high,  
  event_level="first"  
)
```

### 14.2.4 Specificity


- Note that you may have to change `event_level` between “first” and “second” depending on what counts as “success” in your model.

```
logpreds %>% spec(  
  .pred_class,  
  truth = sales_high,  
  event_level="first"  
)
```

### 14.2.5 ROC and AUC

- Note that there are 4 possible ROC curves given by the options
  - `.pred_high` and `.pred_low` (These names will depend on your data.)
  - `event_level="first"` and `event_level="second"`
- You will need to experiment to make sure you choose the right combination for your data.

```
logpreds %>%  
  roc_curve(  
    .pred_high,  
    truth = sales_high,  
    event_level = "first"  
  ) %>%  
  autoplot()
```



```
logpreds %>%
  roc_auc(
    .pred_high,
    truth = sales_high,
    event_level = "first"
  )
```

## 15 Data transforms

### 15.1 Standardising the variables

To apply standardisation to our variables (centring and dividing by the standard deviation) we can use the command `preProcess()` in conjunction with `predict()`

```
mpg_preprocess <- preProcess(mpg)
predict(mpg_preprocess, mpg)
```

### 15.2 Box-Cox transform

#### 15.2.1 Finding $\lambda$ value

```
df_bc<-BoxCoxTrans(y=df$y, x=df$x)
df_bc$lambda
```

The default range for  $\lambda$  is  $[-2, 2]$ . If you want to search over a bigger range, then you can use the `seq()` command with `lambda` option.

```
df_bc<-BoxCoxTrans(y=df$y, x=df$x, lambda=seq(-5,5,0.05))
df_bc$lambda
```

### 15.2.2 Transforming the data

```
predict(df_bc,df$y)
```

## 16 Manipulating time

- Calculate the duration between two date objects in days (as a `difftime` data type)

```
dmy("05-11-2028") - dmy("05-10-2025")
```

- Calculate the duration between two date objects in seconds (as a `duration` data type)

```
as.duration(dmy("05-11-2028") - dmy("05-10-2025"))
```

- If you want these values as integers then wrap the commands in `as.integer()`

```
as.integer(dmy("05-11-2028") - dmy("05-10-2025"))
as.integer(as.duration(dmy("05-11-2028") - dmy("05-10-2025")))
```

## 17 Cleaning data

### 17.1 Finding duplicates

```
df<- tibble(
  x=c(1:7,5,2),
  y=x^2
)
duplicated(df)
```

(To use this to extract the duplicated values, see Section “Extracting rows that match a TRUE/FALSE condition”.)

## 18 Tidying data

- Convert to “long” form (“melting” your data)
  - `gather(dataset, key= "first_new_column", value = "second_new_column", start:end)`

```
data("table4a")
table4a
TB_cases<- gather(table4a, key = "year", value = "cases", `1999`:`2000`)
```

- Convert back to “wide” form, where
  - `key` is the variable that you want to put as new column headings
  - `value` is the variable that you want to put in these new columns
  - `spread(dataset, key= "key_column", value = "data_column")`


```
spread(TB_cases, key = "year", value = "cases")
```



## 19 Plotting


- Scatter plot

```
ggplot(mpg, aes(x=displ, y=cty))+  
  geom_point()
```




- Line plot

```
ggplot(mpg, aes(x=displ, y=cty))+  
  geom_line()
```




- Box plot

```
ggplot(mpg, aes(class, cty))+  
  geom_boxplot()
```




- Bar graph

```
ggplot(mpg, aes(x = manufacturer)) + geom_bar()
```




This is similar to a histogram (but we should use bar charts for categorical variables, and histograms for quantitative variables)

```
ggplot(mpg, aes(manufacturer)) +  
  geom_histogram(stat="count")
```



### 19.1 Rotating x-axis label


```
ggplot(mpg, aes(x = manufacturer)) +  
  geom_bar() +  
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5, hjust=1))
```



## 19.2 Fitting lines and curves to the data


- A straight trend line. Need the method “lm” (as in “linear model”)

```
ggplot(mpg, aes(x=displ, y=cty))+  
  geom_point()+  
  geom_smooth(method="lm")
```



- A possibly curved smoothed average line

```
ggplot(mpg, aes(x=displ, y=cty))+  
  geom_point()+  
  geom_smooth()
```



## 20 R Markdown

### 20.1 Writing equations

- R Markdown uses Latex conventions for equations (so if you know Latex, you can just type it directly into R Markdown)
- To write in-line maths (expressions in the middle of the text) use a single dollar \$ followed by your maths then a second single dollar \$. Eg. this equation  $y = mx + c$  is an in-line maths expression.
  - Note that  $\backslash(\dots\backslash)$  is equivalent to  $\$...\$$  and you can use either.
  - Note that this means you can't just type a dollar symbol in R Markdown. To type a dollar symbol use  $\backslash\$$ . Eg. a \$10 note is blue.
- To write “displayed” maths (expressions separated from the text) use a double dollar \$\$ followed by your maths, then another double dollar \$\$\$. Eg. this equation

$$y = x^3 + 4x + 1$$

is a “displayed” expression.

- Note that  $\backslash[\dots\backslash]$  is equivalent to  $$$\dots$$$$  and you can use either.

### 20.2 Writing long equations

- We can use the `align` environment.
- You end a row in the equation with  $\backslash\backslash$
- Use an ampersand (&) to align the rows



- This also puts an equation number on every row of our equation

$$y = \alpha x^3 + 4\beta x + 3\gamma \tag{1}$$

$$+ 321 - \omega \tag{2}$$

$$= 789 \tag{3}$$

- If you don't want the equation number, use the `align*` environment

$$y = \alpha x^3 + 4\beta x + 3\gamma$$

$$+ 321 - \omega$$

$$= 789$$

## 20.3 Writing maths symbols (including Greek letters)

- The easiest way to write maths symbols, including Greek letters, is to write them as in-line maths. Eg. here is an  $\alpha$ , here is a  $\beta$ , here is a  $\gamma$  and here is an  $\epsilon$ .
  - To type a Greek letter, type a backslash followed by the name of the letter, eg. `\alpha`, `\beta`, `\gamma`, `\epsilon`.
- To put a “hat” on a maths symbol write `\hat{<symbol>}`. Eg. here is a regular  $\alpha$  and here is alpha-hat  $\hat{\alpha}$  using `\hat{\alpha}`.
- To write subscripts use an underscore `{}` after the symbol `<symbol>_{<subscript>}`. Eg. here is a regular  $\alpha$  and here is  $\alpha_{1,2}$  using `\alpha_{1,2}`
  - If you've only got a single character in the subscript, then you don't need the curly brackets `{}`, but it's good practice to use them anyway. Eg. here is a regular  $\alpha$  and here is  $\alpha_1$  (no curly brackets) and here is  $\alpha_1$  (with curly brackets).

### 20.3.1 Some other symbols

- Approximation: `\approx`, ie.  $\approx$

## 20.4 Including R output in your text

- To include some numerical output from R directly into your knitted text use ``r <variable name>``

```
v_1 <- 10.5324
```

Then we can print the value of the variable 10.5324.

- To include the numerical output in an equation, just do the same thing inside `$`, or `$$` or the `align` environments

Our variable is  $v_1 = 10.5324$ , or we could write

$$y = v_1 \alpha$$

$$= 10.5324 \alpha$$